

A Risk-Sensitive Finite-Time Reachability Problem for Safety of Stochastic Dynamic Systems

Margaret P. Chapman¹, Jonathan P. Lacotte², Donggun Lee³, Kevin Smith⁴, Victoria Cheng⁵,
Jaime Fernandez-Fisac¹, Aviv Tamar¹, Susmit Jha⁶, Claire J. Tomlin¹

Abstract—A classic reachability problem for safety of dynamic systems is to compute the set of initial states from which the state trajectory is guaranteed to stay inside a given constraint set over some time horizon. In this paper, we leverage existing theory of reachability analysis and risk measures to formulate a *risk-sensitive* reachability problem for safety of stochastic dynamic systems under non-adversarial disturbances over a finite time horizon. We provide two key contributions to the reachability literature. First, our formulation quantifies the distance between the boundary of the constraint set and the state trajectory for a stochastic dynamic system. In the literature, Hamilton-Jacobi (HJ) reachability methods quantify this distance for non-deterministic systems subject to adversarial disturbances, while stochastic reachability methods reduce the distance to a binary random variable in order to quantify the probability of safety. Second, our formulation accounts for rare high-consequence events by posing the optimal control problem in terms of a risk measure, called *Conditional Value-at-Risk* (CVaR). HJ reachability assumes that high-consequence events occur always, which may yield overly conservative solutions in practice, whereas stochastic reachability does not explicitly account for rare high-consequence events, since the optimal control problem is posed in terms of the expectation operator. We define a *risk-sensitive safe set* as the set of initial states from which the risk of extreme constraint violations can be made small via an appropriate control policy, where risk is quantified using CVaR. We show that certain risk-sensitive safe sets enjoy probabilistic safety guarantees. We provide a dynamic programming algorithm to compute under-approximations of risk-sensitive safe sets and prove the correctness of the algorithm for finite probability spaces. Our proof is a novel contribution, as it does not require the assumption of strong duality, which was required in a previous paper. Finally, we demonstrate the utility of risk-sensitive reachability analysis on a numerical example.

I. INTRODUCTION

Reachability analysis is a formal verification method based on optimal control theory that is used to prove safety or performance properties of dynamic systems [1]. A classic reachability problem for safety is to compute the set of initial states from which the state trajectory is guaranteed to stay inside a given constraint set over some time

horizon. This problem was first considered for discrete-time dynamic systems by Bertsekas and Rhodes under the assumption that disturbances are uncertain but belong to known sets [2], [3], [4]. In this context, the problem is solved using a minimax formulation, in which disturbances behave adversarially and safety is described as a binary notion based on set membership [2], [3], [4].¹

In practice, minimax formulations can yield overly conservative solutions, particularly because disturbances are not often adversarial. Most storms do not cause major floods, and most vehicles are not involved in pursuit-evader games. If there are enough observations of the system, one can estimate a probability distribution for the disturbance, and then assess safety properties of the system in a more realistic context.² For stochastic discrete-time dynamic systems, Abate et al. developed an algorithm that computes the set of initial states from which the probability of safety of the state trajectory can be made large by an appropriate control policy [6].³ Summers and Lygeros extended the algorithm of Abate et al. to quantify the probability of safety and performance of the state trajectory, by specifying that the state trajectory should also reach a target set [7].

Both the stochastic reachability methods [6], [7] and the minimax reachability methods [2], [3], [4] for discrete-time dynamic systems describe safety as a binary notion based on set membership. In Abate et al., for example, the probability of safety to be optimized is the expectation of the product (or maximum) of indicator functions, where each indicator encodes the event that the state at a particular time point is inside a given set [6]. The stochastic reachability methods [6], [7] do not generalize to quantify the random distance between the state trajectory and the boundary of the constraint set, since they use indicator functions to convert probabilities to expectations to be optimized.

In contrast, Hamilton-Jacobi (HJ) reachability methods quantify the deterministic analogue of this distance for continuous-time systems subject to adversarial disturbances (e.g., see [1], [8], [9], [10]). Quantifying the distance between the state trajectory and the boundary of the constraint set in a non-binary fashion may be important in applications where the boundary is not known exactly, or where mild constraint violations are inevitable, but extreme constraint violations must be avoided.

¹M.C., J.F., A.T., and C.T. are with the Department of Electrical Engineering and Computer Sciences, University of California, Berkeley, USA. chapmanm@berkeley.edu

²J.L. is with the Department of Aeronautics and Astronautics, Stanford University, USA.

³D.L. is with the Department of Mechanical Engineering, University of California, Berkeley, USA.

⁴K.S. is with OptiRTC, Inc. and the Department of Civil and Environmental Engineering, Tufts University, USA.

⁵V.C. is with the Department of Civil and Environmental Engineering, University of California, Berkeley, USA.

⁶S.J. is with SRI International, Menlo Park, California, USA.

¹in ref. [4], see Sec. 3.6.2, “Control within a Target Tube”

²Ref. [5] presents methods for estimating probability distributions.

³Safety of the state trajectory is the event that the state trajectory stays in the constraint set over a finite time horizon.

It is imperative that reachability methods for safety take into account the possibility that rare events can occur with potentially damaging consequences. Reachability methods that assume adversarial disturbances (e.g., [1], [3]) suppose that harmful events can always occur, which may yield solutions with limited practical utility, especially in applications with large uncertainty sets. Stochastic reachability methods [6], [7] do not explicitly account for rare high-consequence events because the optimal control problem is expressed as an expectation.

In contrast, we leverage existing results on *risk measures* to formulate an optimal control problem that explicitly encodes a realistic viewpoint on the possibility of rare high-consequence events: harmful events are likely to occur at some point, but they are unlikely to occur always. A *risk measure* is a function that maps a random variable, Z , representing loss into the real line, according to the risk associated with Z (see [11], Sec. 6.3; see [12], Sec. 2.2). Risk-sensitive optimization is being studied in applied mathematics [13], reinforcement learning [14], [15], [16], and optimal control [17].⁴ Risk-sensitive formulations have the potential to inform practical decision-making that also protects against damaging outcomes [18], where the level of conservatism can be modified as needed.

We use a particular risk measure, called *Conditional Value-at-Risk* (CVaR), in this paper. If Z is a random cost with finite expectation, then the Conditional Value-of-Risk of Z at confidence level $\alpha \in (0, 1)$ is,

$$\text{CVaR}_\alpha[Z] = \min_{t \in \mathbb{R}} \left\{ t + \frac{1}{\alpha} \mathbb{E}[\max\{Z - t, 0\}] \right\}; \quad (1)$$

see [11], Equation 6.22.⁵ Note that $\text{CVaR}_\alpha[Z]$ increases from $\mathbb{E}[Z]$ to $\sup Z$, as α decreases from 1 to 0.⁶ Further, there is a well-established relationship between CVaR and chance constraints that we use to obtain probabilistic safety guarantees. Chow et al. provides tractable methods to compute the CVaR of a cumulative cost incurred by a Markov Decision Process [15] that we also leverage. CVaR has additional desirable properties that are of particular interest to researchers in financial risk management and are summarized in ref. [18].

The key contributions of this paper follow. We formulate a risk-sensitive reachability problem for safety of stochastic dynamic systems under non-adversarial disturbances over a finite time horizon. In particular, our formulation quantifies the non-binary distance between the boundary of the constraint set and the state trajectory for a stochastic dynamic system. This is an extension of stochastic reachability methods (e.g., [6], [7]), which reduce this distance to a binary random variable. Further, in contrast to stochastic

reachability methods, our formulation explicitly accounts for rare high-consequence events by posing the optimal control problem in terms of Conditional Value-at-Risk instead of expectation. This is the first use of risk measures in the reachability literature to our knowledge. In Sec. II, we define the notion of a *risk-sensitive safe set* and formalize the problem statement. Sec. ?? summarizes properties of risk-sensitive safe sets, including their relation to probabilistic safety. Sec. IV provides a dynamic programming algorithm to compute under-approximations of risk-sensitive safe sets. In Sec. V, we provide a numerical example in the context of the design of stormwater infrastructure. Sec. VI provides steps for future work.

II. PROBLEM STATEMENT

We consider a stochastic discrete-time dynamic system over a finite time horizon,⁷

$$x_{k+1} = f(x_k, u_k, w_k), \quad k = 0, 1, \dots, N-1, \quad (2)$$

such that $x_k \in \mathbb{R}^n$ is the state of the system at time k , $u_k \in U$ is the control at time k , and $w_k \in D$ is the random disturbance at time k . U and D are finite sets of real-valued vectors. The dynamics function, $f : \mathbb{R}^n \times U \times D \rightarrow \mathbb{R}^n$, is bounded and Lipschitz continuous. The probability that the disturbance equals $d_j \in D$ at time k is, $\mathbb{P}[w_k = d_j] = p_j$, where $0 \leq p_j \leq 1$ and $\sum_{j=1}^W p_j = 1$.⁸ The only source of randomness in the system is the disturbance. The controls are not random. The initial condition, x_0 , is not random. The states, (x_1, \dots, x_N) , are random because they depend on the random disturbance. The collection of *admissible control policies* is,

$$\Pi := \{(\mu_0, \mu_1, \dots, \mu_{N-1}), \text{ where } \mu_k : \mathbb{R}^n \rightarrow U\}. \quad (3)$$

We are given a constraint set, $\mathcal{K} \subset \mathbb{R}^n$, and the safety criterion that the state of the system should stay inside \mathcal{K} over time. For example, if the system is a pond, then x_k may be the water level of the pond at time k , and $\mathcal{K} := [0, 5\text{ft}]$ indicates that the pond overflows if the water level exceeds 5ft. We quantify the extent of constraint violation/satisfaction using a surface function that characterizes the constraint set. Let $g : \mathbb{R}^n \rightarrow \mathbb{R}$ satisfy,

$$x \in \mathcal{K} \iff g(x) < 0, \quad (4)$$

where we adopt the convention provided by [9] in Equation 2.3. For example, we may choose $g(x) := x - 5$ to characterize $\mathcal{K} := [0, 5\text{ft}]$ on the state space, $\mathbb{R}_+ := [0, \infty)$.

We define a *risk-sensitive safe set* as a set of initial states from which risk of extreme constraint violation over time can be made small using an admissible control policy (3), where risk is quantified by *Conditional Value-at-Risk* (1). Formally,

⁴In risk-sensitive optimization, the risk of a cost is minimized, where risk is quantified using a risk measure. Conversely, in stochastic optimization, we usually minimize the expected value of a cost.

⁵Conditional Value-at-Risk is also called *Average Value-at-Risk*, which is abbreviated as AV@R in [11].

⁶Technically, $\text{CVaR}_\alpha[Z] \rightarrow \text{ess sup } Z$ as $\alpha \rightarrow 0$, where $\text{ess sup } Z$ is the *essential supremum* of Z . Informally, essential supremum is a supremum for random variables.

⁷The system model is a special case of the model given by [4] in Sec. 1.2.

⁸We also assume that w_k is independent of x_k , u_k , and disturbances at any other times.

the risk-sensitive safe set at the confidence level, $\alpha \in (0, 1)$, and the risk level, $r \in \mathbb{R}$, is defined as,

$$\mathcal{S}_\alpha^r := \left\{ x \in \mathbb{R}^n \mid \inf_{\pi \in \Pi} \text{CVaR}_\alpha[Z_x^\pi] < r \right\}, \quad (5a)$$

where

$$Z_x^\pi := \max \{g(x_k) \mid k = 0, \dots, N\}, \quad (5b)$$

such that the state trajectory (x_0, x_1, \dots, x_N) evolves according to the dynamics model (2) with the initial state, $x_0 := x$, under the admissible policy, $\pi \in \Pi$. Note that g characterizes the constraint set, \mathcal{K} , according to (4). The goal of this paper is to compute a family of risk-sensitive safe sets at different levels of confidence, $\alpha \in (0, 1)$, and risk, $r \in \mathbb{R}$.

III. RATIONALE

Computing risk-sensitive safe sets is a well-motivated problem for several reasons. This problem is more general than the stochastic reachability problem that is addressed by Abate et al. [6]. Abate et al. solves for the *maximal probabilistic safe set* at any safety level, $\epsilon \in [0, 1]$,

$$\mathcal{S}^*(\epsilon) = \{x \mid \inf_{\pi \in \Pi} \mathbb{E}[Q_x^\pi] \leq \epsilon\}, \quad (6a)$$

where

$$Q_x^\pi := \max \{1_{\mathcal{K}}(x_k) \mid k = 0, \dots, N\}, \quad (6b)$$

such that the state trajectory (x_0, x_1, \dots, x_N) evolves according to a hybrid dynamics model with the initial state, $x_0 := x$, under the admissible policy, $\pi \in \Pi$, and

$$1_{\mathcal{K}}(x) := \begin{cases} 1 & \text{if } x \notin \mathcal{K} \\ 0 & \text{if } x \in \mathcal{K} \end{cases}; \quad (6c)$$

see [6], Equations 11 and 13. If we choose $\alpha := 1$, $g(x) := 1_{\mathcal{K}}(x) - \frac{1}{2}$, and $r := \epsilon - \frac{1}{2}$ in (5), then we can compute (6). [Do you agree? What about the non-strict inequality in $\mathcal{S}^*(\epsilon)$? What is the difference between the set of Markov policies \mathcal{M}_m in Abate, and our Π ? Abate did his formulation for hybrid systems, but we are not doing this explicitly]

Further, risk-sensitive safe sets have two desirable mathematical properties. The first property is that \mathcal{S}_α^r shrinks as the risk level, r , or the confidence level, α , decrease. Since \mathcal{S}_α^r is an r -sublevel set and CVaR_α increases as α decreases, one can show that,

$$\begin{aligned} \mathcal{S}_{\alpha_2}^{r_2} &\subseteq \mathcal{S}_{\alpha_1}^{r_2} \subseteq \mathcal{S}_{\alpha_1}^{r_1}, \text{ and} \\ \mathcal{S}_{\alpha_2}^{r_2} &\subseteq \mathcal{S}_{\alpha_2}^{r_1} \subseteq \mathcal{S}_{\alpha_1}^{r_1} \end{aligned} \quad (7)$$

hold for any $r_1 \geq r_2$ and $1 > \alpha_1 \geq \alpha_2 > 0$. In other words, as the allowable level of risk of constraint violation (r) decreases, or as the fraction of damaging outcomes that are not fully addressed (α) decreases, \mathcal{S}_α^r encodes a higher degree of safety.

The second property is that the risk-sensitive safe sets at risk level, $r := 0$, enjoy probabilistic safety guarantees.

Lemma 1: If $x \in \mathcal{S}_\alpha^0$, then the probability that the state trajectory initialized at x exits the constraint set can be made strictly less than α by an admissible control policy.

Proof: The proof follows from the fact,⁹

$$\text{CVaR}_\alpha[Z_x^\pi] < 0 \implies \mathbb{P}[Z_x^\pi \geq 0] < \alpha.$$

Further, the event, $Z_x^\pi \geq 0$, is equivalent to the event that there exists a state, x_k , of the associated trajectory that exits the constraint set, since $g(x) \geq 0 \iff x \notin \mathcal{K}$.¹⁰ ■

Remark 1: Lemma 1 indicates that \mathcal{S}_α^0 is a subset of Abate et al.'s *maximal probabilistic safe set* at the safety level, α ; see [6], Equations 9 and 11.

IV. COMPUTATIONAL METHOD

The computation of risk-sensitive safe sets is challenging due to the presence of the maximum (as opposed to a summation) and the use of Conditional Value-at-Risk (as opposed to an expectation). In this paper, we provide under-approximations for risk-sensitive safe sets, and an algorithm to compute these under-approximations. Define $\hat{\mathcal{S}}_\alpha^r$ for the confidence level, $\alpha \in (0, 1)$, and the risk level, $r \in \mathbb{R}$,

$$\hat{\mathcal{S}}_\alpha^r := \left\{ x \in \mathbb{R}^n \mid \inf_{\pi \in \Pi} \text{CVaR}_\alpha[\beta Y_x^\pi] < \beta e^{m \cdot r} \right\}, \quad (8a)$$

where

$$Y_x^\pi := \sum_{k=0}^N e^{m \cdot g(x_k)}, \quad (8b)$$

such that the state trajectory (x_0, x_1, \dots, x_N) evolves according to the dynamics model (2) with the initial state, $x_0 := x$, under the admissible policy, $\pi \in \Pi$; $\beta > 0$ and $m > 0$ are constants.

Lemma 2: $\hat{\mathcal{S}}_\alpha^r$ is a subset of the risk-sensitive set, \mathcal{S}_α^r . Also, the gap between $\hat{\mathcal{S}}_\alpha^r$ and \mathcal{S}_α^r can be made smaller by increasing m .

Proof: The proof relies on two facts. The first fact is,

$$\begin{aligned} \max\{x_1, \dots, x_p\} &\leq \log(e^{x_1} + \dots + e^{x_p}) \\ &\leq \max\{x_1, \dots, x_p\} + \log p, \end{aligned} \quad (9)$$

for any $x \in \mathbb{R}^p$; see [19], Sec. 3.1.5 Examples. Using this fact, one can show the following,

$$\begin{aligned} \max\{y_1, \dots, y_p\} &\leq \frac{1}{m} \log(e^{my_1} + \dots + e^{my_p}) \\ &\leq \max\{y_1, \dots, y_p\} + \frac{\log p}{m}, \end{aligned} \quad (10a)$$

for any $y \in \mathbb{R}^p$, $m > 0$. So, as $m \rightarrow \infty$,

$$\frac{1}{m} \log(e^{my_1} + \dots + e^{my_p}) \rightarrow \max\{y_1, \dots, y_p\}. \quad (10b)$$

The second fact is that Conditional Value-at-Risk is a *coherent risk measure*, so it satisfies useful properties. In particular, CVaR is positively homogeneous, i.e.,

$$\text{CVaR}_\alpha[\lambda Z] = \lambda \text{CVaR}_\alpha[Z],$$

for any $\lambda \geq 0$ and random variable Z , and monotonic, i.e.,

$$\text{CVaR}_\alpha[Y] \leq \text{CVaR}_\alpha[Z],$$

⁹The constraint, $\text{CVaR}_\alpha[Z] \leq 0$, gives a conservative approximation of the chance constraint, $\mathbb{P}[Z > 0] \leq \alpha$, for any random variable Z with finite expectation (see [11], Sec. 6.2.4). We do not show $\text{CVaR}_\alpha[Z] < 0 \implies \mathbb{P}[Z > 0] < \alpha$ for brevity.

¹⁰“Associated trajectory” refers to the trajectory that is initialized at x and evolves under the policy, $\pi \in \Pi$, according to the dynamics model (2).

for any random variables, $Y \leq Z$; see [12], Sec. 2.2. Further, CVaR can be expressed as the supremum expectation over a particular set of probability density functions; see [11], Equations 6.40 and 6.70. Using this property and the fact, $\mathbb{E}[\log(Z)] \leq \log(\mathbb{E}[Z])$, one can show,

$$\text{CVaR}_\alpha[\log(Z)] \leq \log(\text{CVaR}_\alpha[Z]), \quad (11)$$

for any random variable, Z , with finite expectation. By monotonicity, positive homogeneity, (10), and (11),

$$\begin{aligned} \text{CVaR}_\alpha[Z_x^\pi] &\leq \frac{1}{m} \text{CVaR}_\alpha[\log(Y_x^\pi)] \\ &\leq \frac{1}{m} \log(\text{CVaR}_\alpha[Y_x^\pi]). \end{aligned} \quad (12)$$

If $x \in \hat{\mathcal{S}}_\alpha^r$, then $\exists \pi \in \Pi$ such that,¹¹

$$\begin{aligned} \text{CVaR}_\alpha[\beta Y_x^\pi] &< \beta e^{m \cdot r} && \iff \\ \text{CVaR}_\alpha[Y_x^\pi] &< e^{m \cdot r} && \iff \\ \frac{1}{m} \log(\text{CVaR}_\alpha[Y_x^\pi]) &< r && \implies \\ \text{CVaR}_\alpha[Z_x^\pi] &< r, \end{aligned}$$

where the last line comes from (12). So, $x \in \mathcal{S}_\alpha^r$. ■

Remark 2: The parameter, β , is included in (8) to counter numerical issues that may arise if m is chosen very large.

Next, we will provide a dynamic programming algorithm to compute our under-approximation, $\hat{\mathcal{S}}_\alpha^r$, at different levels of confidence and risk. The algorithm is based on the insights provided by Chow et al. [15].

A. Algorithm

The value function that characterizes (8),

$$V_0^*(x, \alpha) := \inf_{\pi \in \Pi} \text{CVaR}_\alpha[\beta Y_x^\pi], \quad (13)$$

encodes the minimum risk of cumulative scaled constraint violation of the state trajectory, starting at the initial state, x , and the initial confidence level, α . Eq. (13) can be computed by considering the evolution of a fully observable *augmented state* that consists of the state of the system (2) and the confidence level, $z := (x, \alpha) \in \mathbb{R}^n \times (0, 1)$. We define the augmented set of admissible control policies for the subproblem starting at time k ,

$$\begin{aligned} \bar{\Pi}_k &:= \{(\mu_k, \mu_{k+1}, \dots, \mu_{N-1}), \mu_i : \mathbb{R}^n \times (0, 1) \rightarrow U\}, \\ k &= 0, \dots, N-1, \end{aligned} \quad (14)$$

where both the state of the system and the confidence level at time t are arguments of μ_t . We will provide an algorithm that explicitly computes the value function that is optimized over the augmented policy space,

$$J_0^*(x, \alpha) := \inf_{\pi \in \bar{\Pi}_0} \text{CVaR}_\alpha[\beta Y_x^\pi]. \quad (15)$$

The algorithm depends on an existing result that provides the next confidence level based on the current confidence

¹¹The minimum is attained because the random variables in this paper have finite expectation.

level, the system's current state, and the system's control policy for the sub-problem starting at the current time.

Lemma 3: Lemma 22 of Pflug and Pichler [20] implies the following CVaR-decomposition for the system (2) subject to the augmented policy, $\pi_k \in \bar{\Pi}_k$,

$$\begin{aligned} \text{CVaR}_\alpha[Z^{\pi_k}|x_k] &= \max_{R \in \mathcal{R}(\alpha, \mathbb{P})} \mathbb{E}\left[R \cdot \text{CVaR}_{\alpha R}[Z^{\pi_k}|x_{k+1}]\right|x_k, \alpha] \\ Z^{\pi_k} &:= \sum_{i=k+1}^N c(x_i) \\ \mathcal{R}(\alpha, \mathbb{P}) &:= \left\{R : D \rightarrow \left[0, \frac{1}{\alpha}\right] \text{ such that } \mathbb{E}[R] = 1\right\}, \end{aligned} \quad (16)$$

where $c : \mathbb{R}^n \rightarrow \mathbb{R}$ is a stage cost.

Proof: The proof relies on the measure-theoretic definition of conditional expectation, and the fact that the next state of the system (2) is given by the current state, the current control action, and the current realization of the disturbance. [please provide more details here](#) ■

Remark 3: The term, $\text{CVaR}_\alpha[Z^{\pi_k}|x_k]$, is the risk of the random cumulative cost, Z^{π_k} , of the state trajectory, (x_k, \dots, x_N) , where x_k is the given initial condition, and $\pi_k \in \bar{\Pi}_k$ is the given policy.

Remark 4: For the system (2),

$$\begin{aligned} \mathbb{E}\left[R \cdot \text{CVaR}_{\alpha R}[Z^{\pi_k}|x_{k+1}]\right|x_k, \alpha] &= \\ \sum_{j=1}^W r_j \cdot \text{CVaR}_{\alpha \cdot r_j}[Z^{\pi_k}|x_{k+1}^j] \cdot \mathbb{P}[w_k = d_j], \end{aligned}$$

and

$$\mathbb{E}[R] = \sum_{j=1}^W r_j \cdot \mathbb{P}[w_k = d_j],$$

such that $r_j = R(d_j) \in [0, \frac{1}{\alpha}]$ is the j^{th} realization of the random variable, $R \in \mathcal{R}(\alpha, \mathbb{P})$, and $x_{k+1}^j = f(x_k, \mu_k(x_k, \alpha), d_j)$ is the j^{th} realization of the state of the system at time $k+1$ under the control, $u_k = \mu_k(x_k, \alpha)$, where $w_k = d_j$ is given.

Remark 5: Chow et al. [15] first interpreted Lemma 22 of Pflug and Pichler [20] in the context of Markov Decision Processes. Lemma 3 restates the result in the context of stochastic dynamic systems.

V. NUMERICAL EXAMPLE

VI. CONCLUSION

the conditional value-at-risk is taken with respect to the probability distribution of (w_0, \dots, w_{T-1}) .

VII. CONCLUSION

-inform the cost-effective design of infrastructure that must withstand rare extreme storms, -possible other applications: to reduce overly conservative error bounds that arise in safe dynamic motion planning (e.g., [8]), and to increase the amount of time that an autonomous vehicle can operate safely while simultaneously optimizing for performance.

ACKNOWLEDGMENT

We thank Mo Chen and Jaime Fisac for discussions. M.C. is supported in part by a NSF Graduate Research Fellowship. This work is supported in part by NSF CPS 1740079.




Fig. 1. An illustration of $\text{CVaR}_\delta(Z) \in \mathbb{R}$, if Z is a continuous random variable. The graph shows the probability density function of Z versus the realizations of Z . The area of the right portion under the curve, shown in red, is $\delta \in (0, 1)$. The area of the left portion under the curve, shown in grey, is $1 - \delta$. $\text{CVaR}_\delta(Z)$ is the expectation of the values along the right portion under the curve, indicated by a yellow circle.

REFERENCES

- [1] S. Bansal, M. Chen, S. Herbert, and C. J. Tomlin, "Hamilton-Jacobi Reachability: A Brief Overview and Recent Advances," *arXiv preprint arXiv:1709.07523*, 2017.
- [2] D. P. Bertsekas, "Control of Uncertain Systems with a Set-Membership Description of the Uncertainty," Ph.D. dissertation, Massachusetts Institute of Technology, 1971.
- [3] D. P. Bertsekas and I. B. Rhodes, "On the Minimax Reachability of Target Sets and Target Tubes," *Automatica*, vol. 7, no. 2, pp. 233–247, 1971.
- [4] D. P. Bertsekas, *Dynamic Programming and Optimal Control*, 4th ed. Athena Scientific, 2017, vol. 1.
- [5] B. W. Silverman, *Density Estimation for Statistics and Data Analysis*. Chapman & Hall, 1998.
- [6] A. Abate, M. Prandini, J. Lygeros, and S. Sastry, "Probabilistic reachability and safety for controlled discrete time stochastic hybrid systems," *Automatica*, vol. 44, no. 11, pp. 2724–2734, 2008.
- [7] S. Summers and J. Lygeros, "Verification of discrete time stochastic hybrid systems: A stochastic reach-avoid decision problem," *Automatica*, vol. 46, no. 12, pp. 1951–1961, 2010.
- [8] S. L. Herbert, M. Chen, S. Han, S. Bansal, J. F. Fisac, and C. J. Tomlin, "FaSTrack: A Modular Framework for Fast and Guaranteed Safe Motion Planning," in *Decision and Control (CDC), 2017 IEEE 56th Annual Conference on*. IEEE, 2017, pp. 1517–1522.
- [9] A. Akametalu, "A learning-based approach to safety for uncertain robotic systems," Ph.D. dissertation, EECS Department, University of California, Berkeley, May 2018. [Online]. Available: <http://www2.eecs.berkeley.edu/Pubs/TechRpts/2018/EECS-2018-41.html>
- [10] I. M. Mitchell and J. A. Templeton, "A Toolbox of Hamilton-Jacobi Solvers for Analysis of Nondeterministic Continuous and Hybrid Systems," in *International Workshop on Hybrid Systems: Computation and Control*. Springer, 2005, pp. 480–494.
- [11] A. Shapiro, D. Dentcheva, and A. Ruszczyński, *Lectures on Stochastic Programming: Modeling and Theory*. Society for Industrial and Applied Mathematics, Mathematical Programming Society, 2009.
- [12] J. Kisiala, "Conditional Value-at-Risk: Theory and Applications," Master's thesis, The School of Mathematics, The University of Edinburgh, August 2015.
- [13] A. Ruszczyński, "Risk-averse dynamic programming for Markov decision processes," *Mathematical Programming*, vol. 125, no. 2, pp. 235–261, 2010.
- [14] T. Osogami, "Robustness and Risk-Sensitivity in Markov Decision Processes," in *Advances in Neural Information Processing Systems*, 2012, pp. 233–241.
- [15] Y. Chow, A. Tamar, S. Mannor, and M. Pavone, "Risk-Sensitive and Robust Decision-Making: a CVaR Optimization Approach," in *Advances in Neural Information Processing Systems*, 2015, pp. 1522–1530.
- [16] L. J. Ratliff and E. Mazumdar, "Risk-sensitive inverse reinforcement learning via gradient methods," *arXiv preprint arXiv:1703.09842*, 2017.
- [17] Y.-L. Chow and M. Pavone, "A Framework for Time-consistent, Risk-Averse Model Predictive Control: Theory and Algorithms," in *American Control Conference*. IEEE, 2014, pp. 4204–4211.
- [18] G. Serraino and S. Uryasev, "Conditional Value-at-Risk (CVaR)," in *Encyclopedia of Operations Research and Management Science*. Springer, 2013, pp. 258–266.
- [19] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge University Press, 2004.
- [20] G. C. Pflug and A. Pichler, "Time-consistent decisions and temporal decomposition of coherent risk functionals," *Mathematics of Operations Research*, vol. 41, no. 2, pp. 682–699, 2016.